
Mathematical and Computational Statistics with a
View Towards Finance and Risk Management

Homework 2

Marco Barcellos - 19-770-205
marcoantonio.barcellosjunior@uzh.ch

Fabian Sandmeier - 16-723-710
fabian.sandmeier@uzh.ch

November 6, 2021

Contents

1	True Expected Shortfall	1
2	Implementing the Bootstrap method for a Student's t	1
2.1	Parametric Bootstrap	1
2.2	Non-parametric Bootstrap	2
2.3	Comparing the non-parametric and the parametric Bootstrap	3
3	Simulating the Simulation	4
3.1	The effects of the shape parameter	5
3.2	Using additional information about the underlying sample distribution	6
4	Appendix	9
4.1	True ES - Code	9
4.2	Parametric and non-parametric Bootstrap - Code	10
4.3	CI Coverage - Code	14

1 True Expected Shortfall

In this report, the Expected Shortfall (ES) figures are always reported as the negative of the tail-conditioned expectation, meaning that higher values depict riskier positions.

The function `ES_formula()` is a short program that analytically calculates the ES for a given Student's t random variable. It simply takes the degrees of freedom, location and scale parameters, plus the desired alpha level. The function restricts the degrees of freedom to be greater than 1, otherwise there will be no convergence of the expectation.

```
1 %% True ES
2 ES = -location + (1 / (df - 1)) * (scale / ESlevel) * tpdf(tinv(
    ESlevel, df), df) * (df + tinv(ESlevel, df)^2);
```

Listing 1: Calculating the ES

This function is called multiple times inside the function `ES_bootstrap()`, for both calculating the true ES of a random sample and estimating the true ES from the parameters given by the MLE in the Parametric Bootstrap.

2 Implementing the Bootstrap method for a Student's t

This section shows how the parametric and the non-parametric Bootstrap method can be implemented and how the method can be used to construct a confidence interval for the ES of a given data sample.

The data sample is generated in Matlab as a random draw of n i.i.d. student's t random variables, where n can be specified accordingly and is set to 250 as a default. The degrees of freedom for the underlying distribution must be specified upon the function call.

```
1 %% Generate the random iid sample:
2 random_sample = scale * trnd(df, [n, 1]) + location;
```

Listing 2: Generating the data sample

2.1 Parametric Bootstrap

First, an estimate of the ES is made from the original random sample using the `mle()` function from Matlab. For the given sample, the parameters *df_hat*, *scale_hat* and *location_hat* are being estimated and passed to the function `ES_formula()`. As a result we get *ES_hat_parametric*.

Next, we'd like to produce a Confidence Interval for the ES. This is implemented with a FOR loop that takes the ML estimates of the original sample and produces a new random sample of i.i.d Student's t random variables. We run the `mle()` function on the new bootstrapped sample, obtain the parameter estimates and compute the corresponding estimate of the ES, using again `ES_formula()`. This procedure is repeated B times, obtaining a bootstrapped distribution of our ES estimate.

```

1  for k = 1 : B
2
3      parametric_bootstrap_sample = trnd(df_hat, [n, 1]) *
        scale_hat + location_hat;
4      mle_output = mle(parametric_bootstrap_sample, '
        Distribution', 'tLocationScale');
5
6      % The MLE output for the degrees of freedom may be
        smaller than 1,
7      % for which we don't have a true ES. In this case, we use
        df = 1.01
8      % instead.
9      df_input = max(mle_output(3), 1.01);
10     ES_param_bootstrap(k, 1) = ES_formula(df_input, ESlevel,
        mle_output(1), mle_output(2));
11
12 end

```

Listing 3: Generate B estimates for the ES, using the parametric bootstrap

Last, we take the 5%- and 95%-quantiles of this distribution which becomes our estimated Confidence Interval of the ES.

2.2 Non-parametric Bootstrap

For the non-parametric method, we need to compute the empirical ES of the original sample. This is done by first taking the empirical VaR at the desired quantile level and averaging over the realizations below that. This becomes our empirical estimate of ES.

In order to obtain a Confidence Interval, we run a FOR loop which takes resamples with replacement out of the original sample. Since in this case we are not making any assumptions about the underlying distribution that originated our sample, we are implicitly treating our original sample as if it was the true population distribution.

```

1  for k = 1 : B
2
3      nonparametric_bootstrap_sample = randsample(random_sample
        , n, true);
4      nonparametric_bootstrap_sample_empirical_VaR = -
        quantile(nonparametric_bootstrap_sample, ESlevel);
5      I =(nonparametric_bootstrap_sample < -
        nonparametric_bootstrap_sample_empirical_VaR);
6      ES_nonparam_bootstrap(k, 1) = - mean(
        nonparametric_bootstrap_sample .* I) / ESlevel;
7
8  end

```

Listing 4: Generate B estimates for the ES, using the non-parametric bootstrap

This procedure is again repeated B times, until we obtain a bootstrapped distribution from which we take the 5%- and 95%-quantiles as the boundaries of our Confidence Interval. The

full code for both the parametric and the non-parametric bootstrap is provided in the appendix.

2.3 Comparing the non-parametric and the parametric Bootstrap

Figure 1 shows the result of one run of the `ES_bootstrap()` function. It's possible to see how the distribution of the parametric method is narrower than the non-parametric version as expected due to the higher level of information built into the parametric method. In this particular case, the true ES lies within the boxplot for both methods.






Figure 1: Boxplots for $ES_{5\%}$ with $df = 4$ $n = 250$ and $B = 10000$.

Figures 2a and 2b show the histograms of the bootstrapped samples for both parametric and non-parametric methods for $n = 250$, $B = 10000$ and $ES_{5\%}$. The difference in dispersion in both plots reflects the increased uncertainty associated with the non-parametric method.



(a) Parametric Bootstrap.



(b) Non-parametric Bootstrap.

Figure 2: $ES_{5\%}$ with $df = 3$ $n = 250$ and $B = 5000$.

Both Figures 2a and 2b display a bell-shaped distribution, but this outcome changes significantly for the non-parametric bootstrap when working with $ES_{1\%}$ instead of $ES_{5\%}$. This is due to the small sample size. For $n = 250$, we have few tail events in our sample. The scarcity of samples deep into the tails translates into more difficult inference for lower ES Levels, particularly for the non-parametric method. Since we resample with replacement in this case, the tail events that could manifest in our bootstrapped samples are limited to the ones we actually observed. Since tail events are rarer, the consequence is that for lower ES levels, the histogram of the non-parametric method loses granularity, whereas the parametric one preserves the bell-shaped appearance. This is best illustrated in Figures 3, 4a and 4b.






Figure 3: Boxplots for $ES_{1\%}$ with $df = 4$, $n = 250$ and $B = 10000$.




(a) Parametric Bootstrap.




(b) Non-parametric Bootstrap.

Figure 4: $ES_{1\%}$ with $df = 4$, $n = 250$ and $B = 10000$.

This issue can be addressed with a larger sample size, since the number of tail events then becomes reasonably large again. Figures 5a and 5b depict the histograms for both methods when $n = 25000$. Unfortunately, in practice, obtaining a higher sample size is not always feasible.



(a) Parametric Bootstrap.



(b) Non-parametric Bootstrap.

Figure 5: $ES_{1\%}$ with $df = 4$, $n = 25000$ and $B = 10000$.

3 Simulating the Simulation

In order to check if our CI does its job, meaning that e.g. a 90% CI captures the true ES 90% of the time, we add another layer of simulation on top of our bootstrap method. The function `coverage_ES()` calls the procedure described above *sim* number of times and records for each run whether the reported Confidence Intervals actually capture the true ES. The

actual coverage is then given by the mean of the output vector. Ideally, the nominal coverage, specified as *confidence_level* in the code, should equal the actual coverage, as obtained by simulating the bootstrap method. If the actual coverage is below the nominal coverage, we haven't captured the true ES often enough, meaning that our CI was too narrow.

In addition, for each run the length of the CI is stored and again its mean is reported. Comparing the average length of the CI gives a sense of which method performs better.

```

1  parfor k = 1 : sim
2
3      [trueES, ~, IC_parametric, ~, IC_nonparametric] =
        ES_bootstrap(df, n, B, ESlevel, location, scale);
4      is_in_interval_par(k, 1) = (trueES >= IC_parametric(1) &&
        trueES <= IC_parametric(2));
5      is_in_interval_nonpar(k, 1) = (trueES >= IC_nonparametric
        (1) && trueES <= IC_nonparametric(2));
6      interval_length_par(k, 1) = IC_parametric(2) -
        IC_parametric(1);
7      interval_length_nonpar(k, 1) = IC_nonparametric(2) -
        IC_nonparametric(1);
8
9  end
10
11  is_in_interval = [is_in_interval_par , is_in_interval_nonpar];
12  interval_length = [interval_length_par , interval_length_nonpar];
13
14  mean(is_in_interval)
15  mean(interval_length)

```

Listing 5: Obtain the actual coverage based on *sim* number of simulations

A particular feature of this function is that we implemented it using a PARFOR loop, which in Matlab allows for faster computation of the simulations by parallelizing the work. In our tests, running the `coverage_ES()` function with $sim = 1000$, $B = 500$ and $n = 250$ with the parallel loop takes about 5 minutes.

3.1 The effects of the shape parameter

We also wanted to verify how the actual CI coverage and CI length change for different degrees of freedom. In order to do this, we looped the `coverage_ES()` function for multiple degrees of freedom ranging from 2 to 10 for an ES level of 5%, and 90% nominal CI. This procedure was computationally intensive and required around three hours of run time.

From the plot below, it is possible to see that the parametric bootstrap performs quite well, achieving actual coverage ratios greater than 85% for all df 's for a nominal 90% CI. It is also possible to see that the CI is biased negatively, as the actual coverage never really reaches 90%.

For the non-parametric case, however, the bias is more significant, with actual coverage ratios never quite reaching 75%. The shape parameter doesn't seem to affect the CI coverage that much, except for very low df 's.

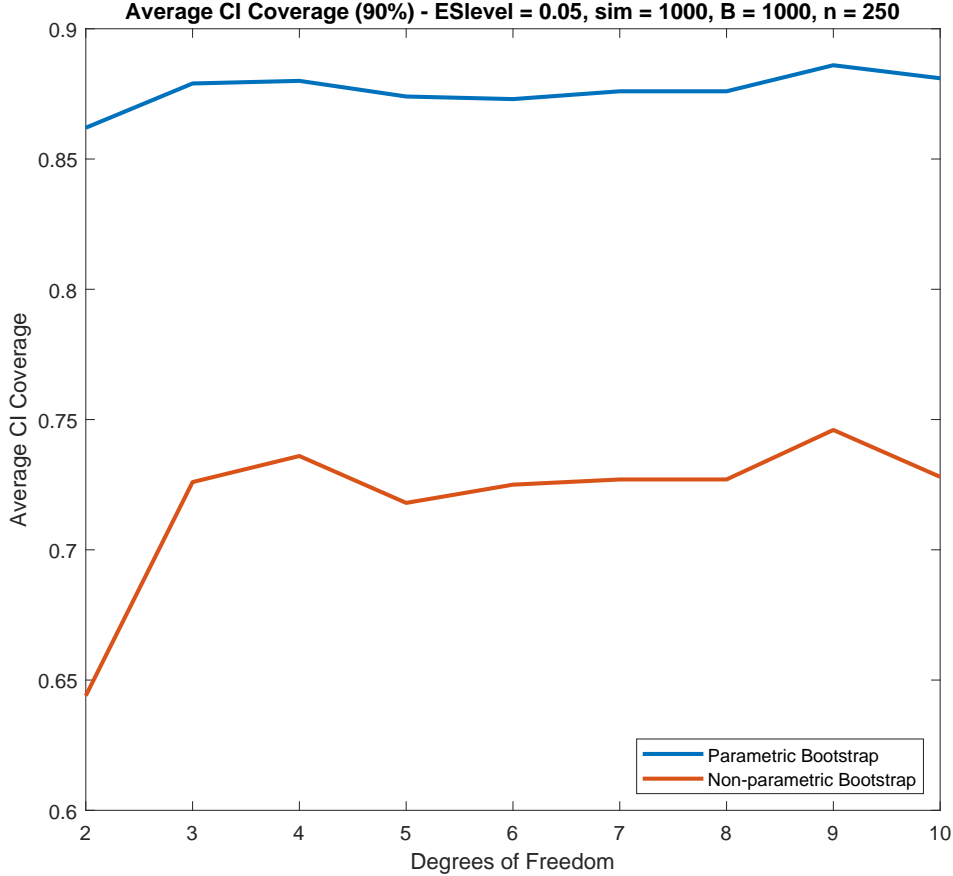


Figure 6: Actual coverage for parametric and non-parametric bootstraps.

Figure 7 shows the average CI length for different degrees of freedom. Here the impact of fatter tails becomes more evident: as degrees of freedom get smaller, the 90% nominal CI becomes larger, reflecting the growing uncertainty about the true ES due to the thicker tails. Once more, the parametric bootstrap performs better than the non-parametric version.

Interestingly, the non-parametric method produced narrower CI for the $df = 2$ case, but, as seen from the low actual coverages in Figure 6, it is clear that the non-parametric method is being too liberal.

Both of these graphics show how much the extra information used by the Parametric Bootstrap matters when calculating CI's. However, this is conditional on correctly assuming the underlying distribution and the performance ranking between the two methods may easily switch under bad assumptions.

3.2 Using additional information about the underlying sample distribution

While in practice we do not have exact details about the parameters of the underlying distribution, let's assume for now that we actually do know some parameters. If we now assume that location and scale parameters are known to be zero and one, respectively, just as the default input arguments of our functions are, we can change the bootstrap procedure to make use of this extra information, and hence expect better performance than the parametric bootstrap method.




Figure 7: Average CI length for parametric and non-parametric bootstraps.

The `mle()` function now fits the data to the PDF of a location 0, scale 1 Student's t and hence forces the location-scale parameters to be the exact ones, therefore only uncertainty about the degrees of freedom remains. This Bootstrap method for obtaining a CI for the ES is again repeated B times.

```

1 %Alternatively, force location=0, scale=1 and only estimate df
2 pd_student_t = @(x,df) tpdf(x,df);
3 paramaters_hat = mle(random_sample, 'pdf', pd_student_t, 'Start'
4     ,1);
5 location_hat = 0;
6 scale_hat = 1;
7 df_hat = paramaters_hat(1);

```

Listing 6: Implment the MLE for location 0, scale 1

```

1 for k = 1 : B
2
3     parametric_bootstrap_sample = trnd(df_hat, [n, 1]) *
4         scale_hat + location_hat;
5     mle_output = mle(parametric_bootstrap_sample, 'pdf',
6         pd_student_t, 'Start', 1);
7
8     % The MLE output for the degrees of freedom may be

```

```

7      smaller than 1,
      % for which we don't have a true ES. In this case, we use
      df = 1.01
8      % instead.
9      df_input = max(mle_output(1), 1.01);
10     ES_param_bootstrap(k, 1) = ES_formula(df_input, ESlevel,
      0, 1);
11
12 end

```

Listing 7: Apply the Bootstrap with location 0, scale 1

Applying this procedure with the same setting as in figure 7 for degrees of freedom 2 and 10, we can see the improvement we get from adding additional information to the bootstrap method. While the actual coverage remained almost the same, since it was already very close to 90% in the first case, the average length of the CI decreases quite a bit.

For 2 df's, we could report an average length of 4.7188 in the case where we made use of the additional information on the parameters, as compared to 5.4282 in the general case. For 10 df's, the average CI length decreased from 0.5743 to 0.5165.

4 Appendix

4.1 True ES - Code

```
1 function ES = ES_formula(df, ESlevel, location, scale)
2
3 %% Verifying inputs are correct and setting up defaults.
4 %(df, ESlevel, location, scale)
5 if nargin == 1
6     ESlevel = 0.05;
7     location = 0;
8     scale = 1;
9 elseif nargin == 2
10    location = 0;
11    scale = 1;
12 elseif nargin == 3
13    scale = 1;
14 end
15
16 %verify inputs
17 if df <= 1
18     msgbox('Degrees of freedom must be greater than 1.')
19     return
20 end
21 if ESlevel >= 1 || ESlevel <= 0
22     msgbox('Alpha must be between 0 and 1.')
23     return
24 end
25
26 %% True ES
27 ES = -location + (1 / (df - 1)) * (scale / ESlevel) * tpdf(
    tinv(ESlevel, df), df) * (df + tinv(ESlevel, df)^2);
28 end
```

Listing 8: Code for ES_formula().

4.2 Parametric and non-parametric Bootstrap - Code

```
1 function [trueES, ES_hat_parametric, IC_parametric,
   ES_hat_nonparametric, IC_nonparametric] = ES_bootstrap(df, n,
   B, ESlevel, location, scale)
2
3 %% TURN OFF MLE WARNING
4 warning('off', 'stats:tlsfit:IterLimit')
5 %warning('on', 'stats:tlsfit:IterLimit')
6 warning('off', 'stats:mlecov:NonPosDefHessian')
7 % warning('on', 'stats:mlecov:NonPosDefHessian')
8
9
10 %% Verifying inputs are correct and setting up defaults.
11 %(df, n, B, ESlevel, location, scale)
12 if nargin == 1
13     n = 250;
14     B = 500;
15     ESlevel = 0.05;
16     location = 0;
17     scale = 1;
18 elseif nargin == 2
19     B = 500;
20     ESlevel = 0.05;
21     location = 0;
22     scale = 1;
23 elseif nargin == 3
24     ESlevel = 0.05;
25     location = 0;
26     scale = 1;
27 elseif nargin == 4
28     location = 0;
29     scale = 1;
30 elseif nargin == 5
31     scale = 1;
32 end
33
34 %verify inputs
35 if df <= 1
36     msgbox('Degrees of freedom must be greater than 1.')
37     return
38 end
39 if ESlevel >= 1 || ESlevel <= 0
40     msgbox('ESlevel must be between 0 and 1.')
41     return
42 end
43 if n <= 0 || B <= 0
44     msgbox('n and B must be greater than zero')
45     return
46 end
47
```

```

48 %% Configuration of confidence interval
49 confidence_level = 0.9;
50 alpha = 1 - confidence_level;
51
52 %% Generate the random iid sample:
53 random_sample = scale * trnd(df, [n, 1]) + location;
54
55 %% True ES by calling the true_ES function:
56 trueES = ES_formula(df, ESlevel, location, scale);
57
58 %% Parametric Bootstrap:
59 % First, uses the random_sample to estimate the parameters of
60 % the
61 % underlying t distribution:
62 paramaters_hat = mle(random_sample, 'Distribution', '
63 tLocationScale');
64 location_hat = paramaters_hat(1);
65 scale_hat = paramaters_hat(2);
66 df_hat = paramaters_hat(3);
67
68 % Second, use the above estimated parameters to calculate the
69 % corresponding ES:
70 % The MLE output for the degrees of freedom may be
71 % smaller than 1,
72 % for which we don't have a true ES. In this case, we use
73 % df = 1.01
74 % instead.
75 df_hat = max(df_hat, 1.01);
76
77 ES_hat_parametric = ES_formula(df_hat, ESlevel, location_hat,
78 scale_hat);
79
80 % Third, FOR loop over B, recreating bootstrap samples out of
81 % the
82 % estimated parameters above and calculating ES via the
83 % parametric formula.
84 ES_param_bootstrap = zeros(B, 1);
85
86 for k = 1 : B
87     parametric_bootstrap_sample = trnd(df_hat, [n, 1]) *
88         scale_hat + location_hat;
89     mle_output = mle(parametric_bootstrap_sample, '
90         Distribution', 'tLocationScale');
91
92     % The MLE output for the degrees of freedom may be
93     % smaller than 1,
94     % for which we don't have a true ES. In this case, we use
95     % df = 1.01
96     % instead.
97     df_input = max(mle_output(3), 1.01);

```

```

87         ES_param_bootstrap(k, 1) = ES_formula(df_input, ESlevel,
88             mle_output(1), mle_output(2));
89     end
90     % Vector ES_param_bootstrap now contains B estimates of ES
91     % calculated with
92     % the parametric formula. In order to obtain a confidence
93     % interval of 90%, we just
94     % need to take the 5%-quantile and the 95%-quantile:
95     IC_parametric = quantile(ES_param_bootstrap, [alpha/2, 1 -
96         alpha/2]);
97
98     % Last, plot a histogram:
99     histogram(ES_param_bootstrap)
100     title('Distribution of estimated ES - Parametric Bootstrap')
101
102     %% Non-parametric Bootstrap:
103     % First, starting from the random sample, calculate the
104     % empirical VaR:
105     empirical_VaR = - quantile(random_sample, ESlevel);
106
107     % Second, calculate the average of the sample tha realized
108     % below the
109     % empirical_VaR:
110     I =(random_sample < - empirical_VaR);
111     ES_hat_nonparametric = - mean(random_sample .* I) / ESlevel;
112
113     % Third, FOR loop where we resample with replacemente from
114     % the original
115     % sample. For each bootstrap sample, calculate the empirical
116     % ES and store
117     % it:
118     ES_nonparam_bootstrap = zeros(B, 1);
119
120     for k = 1 : B
121         nonparametric_bootstrap_sample = randsample(random_sample
122             , n, true);
123         noonparametric_bootstrap_sample_empirical_VaR = -
124             quantile(nonparametric_bootstrap_sample, ESlevel);
125         I =(nonparametric_bootstrap_sample < -
126             noonparametric_bootstrap_sample_empirical_VaR);
127         ES_nonparam_bootstrap(k, 1) = - mean(
128             nonparametric_bootstrap_sample .* I) / ESlevel;
129     end
130
131     % Vector ES_nonparam_bootstrap now contains B estimates of ES
132     % calculated
133     % via quantile of the bootsrtrap resamples. In order to
134     % obtain a confidence
135     % interval of 90%, we just need to take the 5%-quantile and
136     % the 95%-quantile:

```

```

123     IC_nonparametric = quantile(ES_nonparam_bootstrap, [alpha/2,
124         1 - alpha/2]);
125
126     % Last, plot a histogtam
127     figure
128     histogram(ES_nonparam_bootstrap)
129     title('Distribution of estimated ES - Non-Parametric
130         Bootstrap')
131
132     %% Summarizing the two results:
133     figure
134     boxplot([ES_param_bootstrap, ES_nonparam_bootstrap], '
135         orientation', 'horizontal', 'Labels', {'Parametric', 'Non-
136         Parametric'})
137     xline(trueES, 'r--'), legend('True ES')
138
139 end

```

Listing 9: Code for ES_bootstrap().

4.3 CI Coverage - Code

```
1 function [is_in_interval, interval_length] = coverage_ES(sim, df,  
2     n, B, ESlevel, location, scale)  
3  
4     %% Verifying inputs are correct and setting up defaults.  
5     %%(df, n, B, ESlevel, location, scale)  
6     if nargin == 2  
7         n = 250;  
8         B = 500;  
9         ESlevel = 0.05;  
10        location = 0;  
11        scale = 1;  
12    elseif nargin == 3  
13        B = 500;  
14        ESlevel = 0.05;  
15        location = 0;  
16        scale = 1;  
17    elseif nargin == 4  
18        ESlevel = 0.05;  
19        location = 0;  
20        scale = 1;  
21    elseif nargin == 5  
22        location = 0;  
23        scale = 1;  
24    elseif nargin == 6  
25        scale = 1;  
26    end  
27  
28    %verify inputs  
29    if df <= 1  
30        msgbox('Degrees of freedom must be greater than 1.')  
31        return  
32    end  
33    if ESlevel >= 1 || ESlevel <= 0  
34        msgbox('ESlevel must be between 0 and 1.')  
35        return  
36    end  
37    if n <= 0 || B <= 0  
38        msgbox('n and B must be greater than zero')  
39        return  
40    end  
41  
42    % THE CODE BELOW MAKES USE OF PARALLEL COMPUTING  
43  
44    is_in_interval_par = zeros([sim, 1]);  
45    is_in_interval_nonpar = zeros([sim, 1]);  
46    interval_length_par = zeros([sim, 1]);  
47    interval_length_nonpar = zeros([sim, 1]);  
48
```



```

49  parfor k = 1 : sim
50
51      [trueES, ~, IC_parametric, ~, IC_nonparametric] =
          ES_bootstrap(df, n, B, ESlevel, location, scale);
52      is_in_interval_par(k, 1) = (trueES >= IC_parametric(1) &&
          trueES <= IC_parametric(2));
53      is_in_interval_nonpar(k, 1) = (trueES >= IC_nonparametric
          (1) && trueES <= IC_nonparametric(2));
54      interval_length_par(k, 1) = IC_parametric(2) -
          IC_parametric(1);
55      interval_length_nonpar(k, 1) = IC_nonparametric(2) -
          IC_nonparametric(1);
56
57      %loop status
58      k
59
60  end
61
62  is_in_interval = [is_in_interval_par , is_in_interval_nonpar
63      ];
64  interval_length = [interval_length_par ,
65      interval_length_nonpar];
66
67  mean(is_in_interval);
68  mean(interval_length);

```

Listing 10: Code for coverage_ES().